

pa3fzfmwi

November 26, 2024

```
[2]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error, r2_score
from statsmodels.stats.outliers_influence import variance_inflation_factor
import statsmodels.api as sm
```

```
[3]: data = pd.read_csv('Labour_training.csv')
```

```
[4]: data.head()
```

```
[4]:
```

	Age	Education	Race	Hispanic	MaritalStatus	Nodeg	\
0	45	LessThanHighSchool	NotBlack	NotHispanic	Married	1	
1	21	Intermediate	NotBlack	NotHispanic	NotMarried	0	
2	38	HighSchool	NotBlack	NotHispanic	Married	0	
3	48	LessThanHighSchool	NotBlack	NotHispanic	Married	1	
4	18	LessThanHighSchool	NotBlack	NotHispanic	Married	1	

	Earnings_1974	Earnings_1975	Earnings_1978
0	21516.670	25243.550	25564.670
1	3175.971	5852.565	13496.080
2	23039.020	25130.760	25564.670
3	24994.370	25243.550	25564.670
4	1669.295	10727.610	9860.869

```
[5]: data.tail()
```

```
[5]:
```

	Age	Education	Race	Hispanic	MaritalStatus	Nodeg	\
15987	22	HighSchool	black	NotHispanic	NotMarried	0	
15988	20	HighSchool	black	NotHispanic	Married	0	
15989	37	HighSchool	NotBlack	NotHispanic	NotMarried	0	
15990	47	LessThanHighSchool	NotBlack	NotHispanic	Married	1	
15991	40	LessThanHighSchool	NotBlack	NotHispanic	NotMarried	1	

	Earnings_1974	Earnings_1975	Earnings_1978
15987	3975.352	6801.435	2757.438
15988	1445.939	11832.240	6895.072
15989	1733.951	1559.371	4221.865
15990	16914.350	11384.660	13671.930
15991	13628.660	13144.550	7979.724

```
[6]: data.describe()
```

```
[6]:
```

	Age	Nodeg	Earnings_1974	Earnings_1975	Earnings_1978
count	15992.000000	15992.000000	15992.000000	15992.000000	15992.000000
mean	33.225238	0.295835	14016.800304	13650.803376	14846.659673
std	11.045216	0.456432	9569.795893	9270.403225	9647.391524
min	16.000000	0.000000	0.000000	0.000000	0.000000
25%	24.000000	0.000000	4403.452250	4398.823000	5669.298000
50%	31.000000	0.000000	15123.580000	14557.110000	16421.975000
75%	42.000000	1.000000	23584.180000	22923.737500	25564.670000
max	55.000000	1.000000	25862.320000	25243.550000	25564.670000

```
[7]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 15992 entries, 0 to 15991
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Age              15992 non-null  int64
1   Education        15992 non-null  object
2   Race             15992 non-null  object
3   Hisp             15992 non-null  object
4   MaritalStatus    15992 non-null  object
5   Nodeg            15992 non-null  int64
6   Earnings_1974    15992 non-null  float64
7   Earnings_1975    15992 non-null  float64
8   Earnings_1978    15992 non-null  float64
dtypes: float64(3), int64(2), object(4)
memory usage: 1.1+ MB
```

```
[8]: data.isnull().sum()
data = data.dropna()
```

```
[9]: df_encoded = pd.get_dummies(data, columns=['Education', 'Race', 'Hisp', 'MaritalStatus'], drop_first=True)
```

```
[10]: X = df_encoded.drop(columns=['Earnings_1978'])
y = df_encoded['Earnings_1978']
```

```
[11]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)
```

```
[12]: # Train the Linear Regression model
model = LinearRegression()
model.fit(X_train, y_train)
y_pred = model.predict(X_test)

# Residuals
residuals = y_test - y_pred
```

```
[13]: y_pred = model.predict(X_test)
```

```
[14]: mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
```

```
[15]: print(f"Mean Squared Error: {mse:.2f}")
print(f"R-Squared Score: {r2:.2f}")
```

Mean Squared Error: 48625764.00
R-Squared Score: 0.48


```
[16]: coefficients = pd.DataFrame({'Feature': X.columns, 'Coefficient': model.coef_})
print("Feature Coefficients:")
display(coefficients)
```

Feature Coefficients:

	Feature	Coefficient
0	Age	-107.350492
1	Nodeg	-29.339019
2	Earnings_1974	0.283039
3	Earnings_1975	0.475135
4	Education_Intermediate	274.149983
5	Education_LessThanHighSchool	-29.339019
6	Education_PostGraduate	1059.095110
7	Education_graduate	1033.137733
8	Race_black	-857.181352
9	Hisp_hispanic	-442.428781
10	MaritalStatus_NotMarried	-86.350484

```
[24]: plt.figure(figsize=(8, 5))
plt.scatter(y_test, y_pred, alpha=0.7)
plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()],
↳color='red', linestyle='--')
plt.title("Actual vs Predicted (Linearity Check)")
plt.xlabel("Actual Earnings (1978)")
plt.ylabel("Predicted Earnings (1978)")
```


```
plt.show()
```



```
[25]: dw_test = sm.stats.durbin_watson(residuals)
print(f"Durbin-Watson Test Statistic: {dw_test}")
# Values close to 2 indicate no autocorrelation.
```


Durbin-Watson Test Statistic: 1.9323083671337922

```
[26]: plt.figure(figsize=(8, 5))
sns.scatterplot(x=y_pred, y=residuals)
plt.axhline(y=0, color='red', linestyle='--')
plt.title("Residuals vs Predicted (Homoscedasticity Check)")
plt.xlabel("Predicted Earnings (1978)")
plt.ylabel("Residuals")
plt.show()
```



```
[27]: # Q-Q plot
sm.qqplot(residuals, line='45')
plt.title("Q-Q Plot (Normality Check)")
plt.show()

# Histogram of residuals
plt.figure(figsize=(8, 5))
sns.histplot(residuals, kde=True)
plt.title("Residuals Distribution (Normality Check)")
plt.xlabel("Residuals")
plt.show()
```



```
D:\python\Lib\site-packages\seaborn\_oldcore.py:1119: FutureWarning:
use_inf_as_na option is deprecated and will be removed in a future version.
Convert inf values to NaN before operating instead.
  with pd.option_context('mode.use_inf_as_na', True):
```

